

What your renewal data knew, and whether anyone used it

Prepared for the Board · Period 1 Jul 2024 to 30 Jun 2026

Accounts analysed 2,840 Churn events 612 Billing Annual, invoiced Issued 8 Aug 2026

Sample report. Figures are synthetic and illustrate the deliverable's structure, not any real company's results. Every number a real report contains is computed from the client's own exported history.

1 VERDICT

Over the last 24 months you lost **\$5.14M of ARR** to non-renewal and a further **\$2.21M** to contraction, against a starting book of \$47.2M, for a gross revenue retention of **84.4%**. Of the lost ARR, **44% was already flaggable 90 days before the renewal date**, against a random baseline of 13%.

Because you bill annually, every one of these renewal dates was known months in advance. That makes the second number harder to explain than the first: your median first recorded intervention landed **27 days** before the renewal, and **\$2.11M of lost ARR, 41% of the total, has no recorded intervention at all**. This was not a detection problem. The calendar told you when each of these accounts was up.

\$7.35M	84.4%	44%	27d
ARR lost to non-renewal and contraction	Gross revenue retention	Lost ARR flaggable at 90 days	Median first intervention before renewal

2 DATA SUFFICIENCY

Conclusions in this report are drawn only where the underlying data supports them. Cells below the reportability floor are marked and left unquantified rather than printed with confidence intervals too wide to act on.

SOURCE	COVERAGE	EVENT DATING	VERDICT
Billing (NetSuite invoices)	98.4% of accounts	Explicit, to the day	Sufficient
CRM renewal opportunities	97.9%	Explicit	Sufficient
Support tickets	88.1%	Explicit	Sufficient
Product usage	41.3%	Monthly rollup only	Partial, excluded


Product usage was excluded from feature reconstruction: monthly rollups cannot be reconstructed to a daily as-of state without leaking information from after the prediction date. Including it would have inflated every accuracy figure in this report.

REPORTABILITY FLOOR

612 churn events across a 2-way decomposition and 3 lead-time windows gives 6 cells with a median of **102 events each**, comfortably above the 20-event floor. All cells in this report are reportable. Note the decomposition is 2-way rather than 3-way here, which means fewer cells and more events in each of them, so this report is if anything better powered than a card-billed equivalent.

3 WHAT YOU LOST

Revenue leaves in different ways and each has a different fix, so the first thing worth doing is separating them. In your case one of the three usual categories does not apply, and saying so is more useful than reporting it as zero.



ARR lost by category, 24 months. Percentages are of total ARR lost (\$7.35M). The involuntary row is drawn as an empty outline rather than a zero-length bar, because zero and not measurable are different findings.

WHY INVOLUNTARY CHURN IS BLANK, AND WHY THAT IS NOT THE SAME AS ZERO

Involuntary churn means a customer who intended to stay and left anyway because a payment failed. It only exists where a card is charged automatically. You invoice annually against signed renewals, so there is no failed-payment path for a customer to fall down.

We report this as **not measurable** rather than **\$0**. Zero would be a finding, implying we looked and found none. The truth is that the category does not apply to how you bill. Any vendor reporting a clean \$0 here either has not understood your billing model or is padding the report.


WHAT THAT MEANS FOR WHERE THE MONEY IS

In card-billed businesses roughly a quarter of lost revenue is typically involuntary, and it is the cheapest to recover because those customers were not trying to leave. You do not have that pool. Every dollar in the chart above left on purpose, which means all of your recoverable revenue sits in the renewal conversation rather than in billing operations.

4 WHAT WAS KNOWABLE IN ADVANCE

For each account that did not renew, the model was rebuilt using only information that existed at the time, with no later data and no survivorship. The bars show how much of the lost ARR was already in the flagged set at each horizon, beside what a randomly chosen set of the same size would have caught.

■ Model flagged ■ Random baseline, same set size



Share of lost ARR present in the flagged set at each horizon, against a random set of identical size (20 seeded trials, p95 of 16%, 16% and 17% respectively, so every result clears it). Dollar figures are the flagged ARR at that horizon.

HORIZON	FLAGGED ARR	SHARE	BASELINE	RECALL	PRECISION	EVENTS
90 days	\$2.26M	44%	13%	0.41	0.34	251
60 days	\$3.03M	59%	13%	0.56	0.42	343
30 days	\$3.60M	70%	14%	0.67	0.49	410

READ THIS DIFFERENTLY THAN A CARD-BILLED BUSINESS WOULD

On monthly card billing, a cancellation can arrive on any day and lead time measures how much warning the data gave you. You do not have that problem. Every renewal date in this book was on the calendar a year ahead, so the question is not when you could have known the date. It is whether the risk signal on each account was read while there was still time to act on it.

5 HOW FAST YOU REACTED

Against the 612 accounts that did not renew, this is what the record shows.

59%

Lost accounts with any recorded intervention

27d

Median days from first intervention to renewal date

\$2.11M

Lost ARR with no recorded intervention

Twenty-seven days is inside the window where a renewal decision has usually been made internally and a replacement has often been scoped. Read against section 4, where 44% of this revenue was already flagged at 90 days, the pattern is that the signal existed roughly two months before anyone acted on it.

READ THIS BEFORE ACTING ON THE NUMBER ABOVE

This measures *recorded* intervention. Renewal conversations held by phone, over email outside the CRM, or in a QBR without a logged follow-up are invisible here and will inflate the no-intervention figure. Treat \$2.11M as an upper bound on unattended ARR rather than a settled fact. If your team believes the real figure is materially lower, that gap is itself worth knowing, because it means renewal activity is not being captured anywhere a successor could find it.

6 MODEL TRACK RECORD

Ten walk-forward folds, each trained only on data preceding its origin date, with a 14-day embargo between training and scoring windows to prevent boundary leakage.

METRIC	RESULTINTERPRETATION
Precision @ 20	0.4995% CI 0.42 to 0.56
Lift over base rate	3.6xBase rate 13.6%
ROC AUC (median fold)	0.74Range 0.69 to 0.78 across 10 folds
Separable from chance	YesCI excludes the baseline at every horizon
Negative control	0.503Shuffled labels, no residual signal
Leakage canary	CleanNo feature exceeds its as-of availability

The negative control and leakage canary are reported because their failure would invalidate everything above them. A model that scores 0.497 on shuffled labels is finding structure in the data rather than in the evaluation.

7 WHAT THIS REPORT DOES NOT PROVE

Everything above is **prediction and correlation**. Nothing in this report establishes that intervening on a flagged account would have prevented its cancellation.

Specifically:

- **The reaction-time findings are correlational.** Accounts that received early contact may differ systematically from those that did not, being larger, better covered or more visible. That selection effect is not controlled here.
- **Flagged-and-saved is not evidence of a save.** Some flagged accounts would have renewed regardless. Without a control arm, that fraction is unknown.
- **Historical accuracy is not a forward guarantee.** It measures a model against a period that has already happened, under conditions that may not persist.

Establishing causation requires a randomised holdout running forward: a stratified share of flagged accounts deliberately withheld from outreach, assigned before outcomes are known. That is the only design that answers "did the intervention work," and this report is not a substitute for it.

8 SEALED FORWARD FORECAST

The forecast below was committed before the period began, and its hash was published at the time of issue. It cannot be revised after the fact. It will be scored against outcomes on **31 December 2026** and the result published whether or not it is favourable.

COMMITMENT	FORECAST	INTERVAL
Lost ARR, Q3 to Q4 2026	\$2.41M	\$2.06M to \$2.79M
Gross revenue retention	85.2%	83.4% to 86.8%
Share flaggable at 90 days	45%	38% to 52%
Accounts named in the sealed at-risk list	214	,

Commitment hash (SHA-256):
b71e4c05fa39d8672c1ab5e04938f7d6250ce8a1477b93fc602d5e8143a9017c

Verification: the sealed list and its hash were recorded on issue. Re-hashing the delivered list must reproduce the value above. A mismatch means the forecast was altered and the result should be discarded.

Method. Account state was reconstructed to its as-of value at each prediction date from an append-only event ledger. Labels were derived from billing and subscription events, never from present-day account status. Evaluation is walk-forward with out-of-time splits, a 14-day embargo, and isotonic calibration fitted within each training window only.

Sample report · synthetic figures · prepared 8 August 2026 · Indexxero